

# NBAPredictionCode

December 19, 2018

```
In [3]: # Import necessary packages
```

```
%matplotlib inline

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.neural_network import MLPClassifier
import sklearn.metrics as metrics
from sklearn.model_selection import KFold
from operator import itemgetter
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import RandomizedSearchCV
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn import neighbors
import seaborn as sns
from sklearn import linear_model, preprocessing
```

```
In [4]: # Import the datasets
```

```
rookieDf = pd.read_csv('2017-rookies.csv')
pastDf = pd.read_csv('historical-rookies.csv')
```

```
# Preview the historical dataset
```

```
pastDf.head()
```

```
Out[4]:
```

	SeasonID	Player	Season	Age	Tm	Lg	Pos	G	GS	MPG	\
0	0	Alaa Abdelnaby	1990-91	22	POR	NBA	PF	43	0	6.7	
1	43	Mahmoud Abdul-Rauf	1990-91	21	DEN	NBA	PG	67	19	22.5	
2	52	Tariq Abdul-Wahad	1997-98	23	SAC	NBA	SG	59	16	16.3	
3	62	Shareef Abdur-Rahim	1996-97	20	VAN	NBA	PF	80	71	35.0	
4	88	Alex Abrines	2016-17	23	OKC	NBA	SG	66	6	15.5	
	...	DWS	WS	WS/48	OBPM	DBPM	BPM	VORP	PlayerID	Hall of Fame	\
0	...	0.5	0.5	0.079	-4.2	-0.7	-5.0	-0.2	0	0	

1	...	-0.3	-1.0	-0.031	-1.7	-4.4	-6.1	-1.6	3	0
2	...	0.6	-0.2	-0.008	-4.2	-1.7	-5.9	-0.9	4	0
3	...	1.2	2.9	0.049	-0.8	-1.2	-2.0	0.0	5	0
4	...	0.9	2.0	0.094	-0.4	-2.3	-2.7	-0.2	9	0

	All Star
0	0
1	0
2	0
3	1
4	0

[5 rows x 55 columns]

In [5]: `pastDf.columns.values`

Out[5]: `array(['SeasonID', 'Player', 'Season', 'Age', 'Tm', 'Lg', 'Pos', 'G', 'GS', 'MPG', 'FG/G', 'FGA/G', 'FG%', '3P/G', '3PA/G', '3P%', '2P/G', '2PA/G', '2P%', 'eFG%', 'FT/G', 'FTA/G', 'FT%', 'ORB/G', 'DRB/G', 'TRB/G', 'AST/G', 'STL/G', 'BLK/G', 'TOV/G', 'PF/G', 'PTS/G', 'PER', 'TS%', '3PAr', 'FTr', 'ORB%', 'DRB%', 'TRB%', 'AST%', 'STL%', 'BLK%', 'TOV%', 'USG%', 'OWS', 'DWS', 'WS', 'WS/48', 'OBPM', 'DBPM', 'BPM', 'VORP', 'PlayerID', 'Hall of Fame', 'All Star'], dtype=object)`

In [6]: *# Preview the rookies dataset*

`rookieDf.head()`

Out[6]:	SeasonID	Player	Season	Tm	Lg	Pos	G	GS	MPG	FG/G	\
0	0	Ben Simmons	2017-2018	PHI	NBA	PG	81	81	33.7	6.7	
1	1	Lonzo Ball	2017-18	LAL	NBA	PG	52	50	34.2	3.9	
2	2	Jayson Tatum	2017-18	BOS	NBA	SF	80	80	30.5	5.0	
3	3	Josh Jackson	2017-18	PHO	NBA	SF	77	35	25.4	5.1	
4	4	De'Aaron Fox	2017-18	SAC	NBA	PG	73	60	27.8	4.5	
	...	USG%	OWS	DWS	WS	WS/48	OBPM	DBPM	BPM	VORP	PlayerID
0	...	22.3	4.2	5.0	9.2	0.162	1.0	3.6	4.6	4.6	0
1	...	17.4	-0.5	2.5	2.0	0.053	-0.8	2.5	1.7	1.7	1
2	...	19.5	3.0	4.0	7.1	0.139	-0.5	1.5	1.0	1.8	2
3	...	26.0	-2.0	1.2	-0.7	-0.018	-3.1	-1.2	-4.3	-1.1	3
4	...	23.4	-1.7	1.1	-0.6	-0.014	-2.8	-1.5	-4.4	-1.2	4

[5 rows x 52 columns]

In [18]: `pastDf.dtypes`

Out[18]:	SeasonID	int64
	Player	object

Season	object
Age	int64
Tm	object
Lg	object
Pos	object
G	int64
GS	int64
MPG	float64
FG/G	float64
FGA/G	float64
FG%	float64
3P/G	float64
3PA/G	float64
3P%	float64
2P/G	float64
2PA/G	float64
2P%	float64
eFG%	float64
FT/G	float64
FTA/G	float64
FT%	float64
ORB/G	float64
DRB/G	float64
TRB/G	float64
AST/G	float64
STL/G	float64
BLK/G	float64
TOV/G	float64
PF/G	float64
PTS/G	float64
PER	float64
TS%	float64
3PAr	float64
FTr	float64
ORB%	float64
DRB%	float64
TRB%	float64
AST%	float64
STL%	float64
BLK%	float64
TOV%	float64
USG%	float64
OWS	float64
DWS	float64
WS	float64
WS/48	float64
OBPM	float64
DBPM	float64

```

BPM                float64
VORP               float64
PlayerID           int64
Hall of Fame       int64
All Star           int64
dtype: object

```

## 1 Logistic Regression

```


In [9]: ytrain = np.array(pastDf['All Star'])
        Xtrain = np.array(pastDf[['G', 'MPG', 'FG/G', 'FGA/G', '2P%', '3P%', 'FT%', 'TRB/G', 'AST/G', 'TS%', '3PAr', 'FTr']])

        from sklearn import linear_model

        regr = linear_model.LogisticRegression()
        regr.fit(Xtrain, ytrain)
        ytrain_pred = regr.predict(Xtrain)
        plt.plot(Xtrain, ytrain, 'ro')
        plt.plot(Xtrain, ytrain_pred, 'ro')
        plt.legend(['actual', 'predicted'])
        plt.xlabel('PTS/G')
        plt.ylabel('All Star')

```

Out[9]: Text(0,0.5,'All Star')



## 2 Prepare Rookie Dataset for Predictions

```
In [56]: # Prepare rookies dataframe to be plugged into machine learning models by skipping co
rookieNames = rookieDf.iloc[:, 1]
rookieFeatures = rookieDf[['G', 'MPG', 'FG/G', 'FGA/G', '2P%', '3P%', 'FT%', 'TRB/G',
                           'TS%', '3PAr', 'FTr']]
```

## 3 Cross Validation - Hall of Fame

```
In [45]: vals, y = np.unique(pastDf['Hall of Fame'].values, return_inverse=True)
```

```
xnames = pastDf.columns[9:25]
X= np.array(pastDf[xnames].values)
Xs = preprocessing.scale(X)

logreg= linear_model.LogisticRegression(C=1e5)
logreg.fit(Xs, y)
```

```
Out[45]: LogisticRegression(C=100000.0, class_weight=None, dual=False,
                             fit_intercept=True, intercept_scaling=1, max_iter=100,
                             multi_class='ovr', n_jobs=1, penalty='l2', random_state=None,
                             solver='liblinear', tol=0.0001, verbose=0, warm_start=False)
```

```
In [46]: from sklearn.model_selection import KFold
         from sklearn.metrics import precision_recall_fscore_support
```

```
nfold = 10
kf = KFold(n_splits = nfold, shuffle=True)
```

```
rec = []
f1 = []
prec = []
err_rate = []
```

```
for Itr, Its in kf.split(Xs):
    Xtr = Xs[Itr,:]
    ytr = y [Itr]
    Xts = Xs[Its,:]
    yts = y[Its]
```

```
logreg.fit(Xtr,ytr)
```

```
yhat = logreg.predict(Xts)
```

```
preci,reci,f1i,_= precision_recall_fscore_support(yts,yhat,average='binary')
```

```
prec.append(preci)
rec.append(reci)
```

```

f1.append(f1i)
err_rate.append(np.mean(yts!=yhat))

prec = np.mean(prec)
rec = np.mean(rec)
f1 = np.mean(f1)
err_mean= np.mean(err_rate)

print('Precision ' + str(prec))
print('Recall ' + str(rec))
print('f1 ' + str(f1))
print('error rate ' + str(err_mean))

```

```

Precision 0.275
Recall 0.09916666666666665
f1 0.13071428571428573
error rate 0.02140475563744845

```

```

/anaconda3/lib/python3.6/site-packages/sklearn/metrics/classification.py:1135: UndefinedMetric
'precision', 'predicted', average, warn_for)
/anaconda3/lib/python3.6/site-packages/sklearn/metrics/classification.py:1135: UndefinedMetric
'precision', 'predicted', average, warn_for)

```

## 4 Make Prediction - Hall of Fame

```

In [58]: ynew = logreg.predict(rookieFeatures)
         for i in range(len(rookieFeatures)):
             print("Name = %s, Predicted = %s" % (rookieNames[i], ynew[i]))

```

```

Name = Ben Simmons, Predicted = 1
Name = Lonzo Ball, Predicted = 0
Name = Jayson Tatum, Predicted = 1
Name = Josh Jackson, Predicted = 1
Name = De'Aaron Fox, Predicted = 1
Name = Lauri Markkanen, Predicted = 1
Name = Frank Ntilikina, Predicted = 0
Name = Dennis Smith, Predicted = 1
Name = Zach Collins, Predicted = 1
Name = Malik Monk, Predicted = 1
Name = Luke Kennard, Predicted = 1
Name = Donovan Mitchell, Predicted = 1
Name = Bam Adebayo, Predicted = 1
Name = Justin Jackson, Predicted = 1
Name = T.J. Leaf, Predicted = 1
Name = John Collins, Predicted = 1
Name = Terrance Ferguson, Predicted = 1

```

```

Name = Jarrett Allen, Predicted = 1
Name = OG Anunoby, Predicted = 1
Name = Kyle Kuzma, Predicted = 1
Name = Josh Hart, Predicted = 1
Name = Wesley Iwundu, Predicted = 1
Name = Frank Mason, Predicted = 1
Name = Semi Ojeleye, Predicted = 1
Name = Jordan Bell, Predicted = 1
Name = Dwayne Bacon, Predicted = 1
Name = Tyler Dorsey, Predicted = 1
Name = Dillon Brooks, Predicted = 1
Name = Sterling Brown, Predicted = 1
Name = Sindarius Thornwell, Predicted = 1

```

## 5 Cross Validation - All Stars

```

In [59]: vals, y = np.unique(pastDf['All Star'].values, return_inverse=True)

```

```

xnames = pastDf.columns[9:25]
X= np.array(pastDf[xnames].values)
Xs = preprocessing.scale(X)

```

```

logreg= linear_model.LogisticRegression(C=1e5)
logreg.fit(Xs, y)

```

```

Out[59]: LogisticRegression(C=100000.0, class_weight=None, dual=False,
    fit_intercept=True, intercept_scaling=1, max_iter=100,
    multi_class='ovr', n_jobs=1, penalty='l2', random_state=None,
    solver='liblinear', tol=0.0001, verbose=0, warm_start=False)

```

```

In [60]: nfold = 10
kf = KFold(n_splits = nfold, shuffle=True)

```

```

rec = []
f1 = []
prec = []
err_rate = []

```

```

for Itr, Its in kf.split(Xs):
    Xtr = Xs[Itr,:]
    ytr = y [Itr]
    Xts = Xs[Its,:]
    yts = y[Its]

    logreg.fit(Xtr,ytr)

    yhat = logreg.predict(Xts)

```

```

    preci, reci, f1i, _ = precision_recall_fscore_support(yts, yhat, average='binary')

    prec.append(preci)
    rec.append(reci)
    f1.append(f1i)
    err_rate.append(np.mean(yts!=yhat))

prec = np.mean(prec)
rec = np.mean(rec)
f1 = np.mean(f1)
err_mean = np.mean(err_rate)

print('Precision ' + str(prec))
print('Recall ' + str(rec))
print('f1 ' + str(f1))
print('error rate ' + str(err_mean))

```

```

Precision 0.6427114552114552
Recall 0.2845761950940979
f1 0.39168541010646274
error rate 0.0748910532011348

```

## 6 Make Prediction - All Star

```

In [62]: ynew = logreg.predict(rookieFeatures)
        for i in range(len(rookieFeatures)):
            print("Name = %s, Predicted = %s" % (rookieNames[i], ynew[i]))

```

```

Name = Ben Simmons, Predicted = 1
Name = Lonzo Ball, Predicted = 1
Name = Jayson Tatum, Predicted = 1
Name = Josh Jackson, Predicted = 1
Name = De'Aaron Fox, Predicted = 1
Name = Lauri Markkanen, Predicted = 1
Name = Frank Ntilikina, Predicted = 1
Name = Dennis Smith, Predicted = 1
Name = Zach Collins, Predicted = 1
Name = Malik Monk, Predicted = 1
Name = Luke Kennard, Predicted = 1
Name = Donovan Mitchell, Predicted = 1
Name = Bam Adebayo, Predicted = 1
Name = Justin Jackson, Predicted = 1
Name = T.J. Leaf, Predicted = 1
Name = John Collins, Predicted = 1
Name = Terrance Ferguson, Predicted = 1
Name = Jarrett Allen, Predicted = 1

```



```

Name = OG Anunoby, Predicted = 1
Name = Kyle Kuzma, Predicted = 1
Name = Josh Hart, Predicted = 1
Name = Wesley Iwundu, Predicted = 1
Name = Frank Mason, Predicted = 1
Name = Semi Ojeleye, Predicted = 1
Name = Jordan Bell, Predicted = 1
Name = Dwayne Bacon, Predicted = 1
Name = Tyler Dorsey, Predicted = 1
Name = Dillon Brooks, Predicted = 1
Name = Sterling Brown, Predicted = 1
Name = Sindarius Thornwell, Predicted = 1

```

## 7 Creating data splits for HOF models

```

In [63]: trainData, testData = train_test_split(pastDf, test_size = 0.25, random_state = 0)

xtrainData = trainData[['G', 'MPG', 'FG/G', 'FGA/G', '2P%', '3P%', 'FT%', 'TRB/G', 'AST',
                        'TS%', '3PAr', 'FTr']]

xtestData = testData[['G', 'MPG', 'FG/G', 'FGA/G', '2P%', '3P%', 'FT%', 'TRB/G', 'AST',
                      'TS%', '3PAr', 'FTr']]

ytrainData = trainData[['Hall of Fame']]

ytestData = testData[['Hall of Fame']]

```

## 8 Create models and confusion matrices for HOF models

```

In [66]: svcModel = SVC(kernel='rbf', gamma=1e-4, C=10, probability = True)

svcModel.fit(xtrainData, ytrainData.values.ravel())

ysvc = svcModel.predict(xtestData)

probability = svcModel.predict_proba(xtestData)

posProbability = probability[:, 1]

fprSVC, tprSVC, thresholdSVC = metrics.roc_curve(ytestData, posProbability)
roc_aucSVC = metrics.auc(fprSVC, tprSVC)

cvScoreSVC = cross_val_score(svcModel, xtestData, ytestData.values.ravel(), cv = 3, s

```

```

print("Accuracy score: %.3f" % metrics.accuracy_score(ytestData, ysvc))
print("Log loss: %.3f" % metrics.log_loss(ytestData, probability))
print("Area under ROC curve: %.3f" % metrics.roc_auc_score(ytestData, posProbability))
print("Accuracy (cross validation score): %0.2f (+/- %0.2f)" % (cvScoreSVC.mean(), cvScoreSVC.std()))

```

Accuracy score: 0.983

Log loss: 0.075

Area under ROC curve: 0.951

Accuracy (cross validation score): 0.98 (+/- 0.00)

```
In [68]: confusionmatrix = metrics.confusion_matrix(ytestData, ysvc)
```

```
svcHofconfusionmatrix, ax = plt.subplots()
```

```
sns.heatmap(confusionmatrix, annot=True, ax = ax, linewidth = 2)
```

```
ax.set_xlabel("Predicted")
```

```
ax.set_ylabel("Actual")
```

```
labels = ["Not HOF", "HOF"]
```


```
ax.set_xticklabels(labels)
```

```
ax.set_yticklabels(labels)
```

```
svcHofconfusionmatrix.suptitle("SVC Confusion Matrix", weight = 'bold', size = 18, y = 1.05)
```

```
svcHofconfusionmatrix.savefig('svc-hof-cm.png', dpi = 400, bbox_inches = 'tight')
```

## SVC Confusion Matrix



```
In [69]: deepneuralnet = MLPClassifier(
    solver='lbfgs',
    hidden_layer_sizes=100,
    max_iter=10000,
    shuffle=False,
    random_state=0,
    activation='identity')

deepneuralnet.fit(xtrainData, ytrainData.values.ravel())

y_deepneuralnet = deepneuralnet.predict(xtestData)

print("Accuracy score: %.3f" % metrics.accuracy_score(ytestData, y_deepneuralnet))

proba = deepneuralnet.predict_proba(xtestData)
print("Log loss: %.3f" % metrics.log_loss(ytestData, proba))

posProb = proba[:, 1]
print("Area under ROC curve: %.3f" % metrics.roc_auc_score(ytestData, posProb))

fprDNN, tprDNN, thresholdDNN = metrics.roc_curve(ytestData, posProb)
```

```

roc_aucDNN = metrics.auc(fprDNN, tprDNN)

cvScoreDNN = cross_val_score(deepneuralnet, xtestData, ytestData.values.ravel(), cv =
print("Accuracy (cross validation score): %0.2f (+/- %0.2f)" % (cvScoreDNN.mean(), cv

Accuracy score: 0.988
Log loss: 0.045
Area under ROC curve: 0.971
Accuracy (cross validation score): 0.98 (+/- 0.00)

In [71]: confusionmatrix = metrics.confusion_matrix(ytestData, y_deepneuralnet)

deepneuralnetHofCM, ax = plt.subplots()

sns.heatmap(confusionmatrix, annot=True, ax = ax, linewidth = 2)

ax.set_xlabel("Predicted")
ax.set_ylabel("Actual")


labels = ["Not HOF", "HOF"]
ax.set_xticklabels(labels)
ax.set_yticklabels(labels)

deepneuralnetHofCM.suptitle("DNN Confusion Matrix", weight = 'bold', size = 18, y = 1

deepneuralnetHofCM.savefig('dnn-hof-cm.png', dpi = 400, bbox_inches = 'tight')

```

## DNN Confusion Matrix



## 9 Randomized search CV

In [73]: # SVC

```
C = [int(x) for x in np.linspace(start = 1, stop = 100, num = 10)]

kernel = ['rbf', 'linear']

gamma = [float(x) for x in np.linspace(start = 1e-5, stop = 10, num = 10)]

random_grid = {'C': C,
               'kernel': kernel,
               'gamma': gamma}

svc_random = RandomizedSearchCV(estimator = svcModel, param_distributions = random_grid,
                                verbose=2, random_state=42, n_jobs = 20)
```

In [74]: `svc_random.fit(xtrainData, ytrainData.values.ravel())`

```
y_svcrand = svc_random.predict(xtestData)
```

```

searchScore = metrics.accuracy_score(ytestData, y_svcrand)
nonSearchScore = metrics.accuracy_score(ytestData, y_svc)
improvement = (searchScore - nonSearchScore) / nonSearchScore

```

```

print("Percent improvement: %.5f" % improvement)

```

Fitting 3 folds for each of 25 candidates, totalling 75 fits

```

[CV] kernel=linear, gamma=7.77778, C=45 ...
[CV] kernel=linear, gamma=7.77778, C=45 ...
[CV] kernel=linear, gamma=7.77778, C=45 ...
[CV] kernel=linear, gamma=7.77778, C=1 ...
[CV] kernel=linear, gamma=7.77778, C=1 ...
[CV] kernel=linear, gamma=7.77778, C=1 ...
[CV] kernel=rbf, gamma=5.55556, C=12 ...
[CV] kernel=rbf, gamma=5.55556, C=12 ...
[CV] kernel=rbf, gamma=5.55556, C=12 ...
[CV] kernel=rbf, gamma=10.0, C=78 ...
[CV] kernel=rbf, gamma=10.0, C=78 ...
[CV] kernel=rbf, gamma=10.0, C=78 ...
[CV] kernel=rbf, gamma=4.44445, C=67 ...
[CV] kernel=rbf, gamma=4.44445, C=67 ...
[CV] kernel=linear, gamma=7.77778, C=56 ...
[CV] kernel=rbf, gamma=4.44445, C=67 ...
[CV] kernel=linear, gamma=7.77778, C=56 ...
[CV] kernel=linear, gamma=7.77778, C=56 ...
[CV] kernel=linear, gamma=4.44445, C=34 ...
[CV] kernel=linear, gamma=4.44445, C=34 ...
[CV] ... kernel=rbf, gamma=10.0, C=78, total= 3.1s
[CV] kernel=linear, gamma=4.44445, C=34 ...
[CV] ... kernel=rbf, gamma=10.0, C=78, total= 3.1s
[CV] kernel=rbf, gamma=5.55556, C=89 ...
[CV] ... kernel=rbf, gamma=10.0, C=78, total= 3.2s
[CV] kernel=rbf, gamma=5.55556, C=89 ...

```

```

[Parallel(n_jobs=20)]: Done 1 tasks | elapsed: 3.5s

```

```

[CV] ... kernel=rbf, gamma=5.55556, C=12, total= 3.3s
[CV] kernel=rbf, gamma=5.55556, C=89 ...
[CV] ... kernel=rbf, gamma=5.55556, C=12, total= 3.3s
[CV] ... kernel=rbf, gamma=4.44445, C=67, total= 3.3s
[CV] ... kernel=rbf, gamma=5.55556, C=12, total= 3.4s
[CV] kernel=rbf, gamma=7.77778, C=89 ...
[CV] kernel=rbf, gamma=7.77778, C=89 ...
[CV] kernel=rbf, gamma=7.77778, C=89 ...
[CV] ... kernel=rbf, gamma=4.44445, C=67, total= 3.3s
[CV] kernel=linear, gamma=2.22223, C=23 ...

```

[CV] ... kernel=rbf, gamma=4.44445, C=67, total= 3.3s  
 [CV] kernel=linear, gamma=2.22223, C=23 ...  
 [CV] ... kernel=rbf, gamma=5.55556, C=89, total= 3.0s  
 [CV] kernel=linear, gamma=2.22223, C=23 ...  
 [CV] ... kernel=rbf, gamma=5.55556, C=89, total= 3.0s  
 [CV] kernel=rbf, gamma=3.33334, C=34 ...  
 [CV] ... kernel=rbf, gamma=7.77778, C=89, total= 2.9s  
 [CV] kernel=rbf, gamma=3.33334, C=34 ...  
 [CV] ... kernel=rbf, gamma=5.55556, C=89, total= 3.0s  
 [CV] kernel=rbf, gamma=3.33334, C=34 ...  
 [CV] ... kernel=rbf, gamma=7.77778, C=89, total= 2.9s  
 [CV] ... kernel=rbf, gamma=7.77778, C=89, total= 2.9s  
 [CV] kernel=rbf, gamma=1.11112, C=100 ...  
 [CV] kernel=rbf, gamma=1.11112, C=100 ...  
 [CV] ... kernel=rbf, gamma=3.33334, C=34, total= 3.1s  
 [CV] ... kernel=rbf, gamma=3.33334, C=34, total= 3.2s  
 [CV] kernel=rbf, gamma=1.11112, C=100 ...  
 [CV] kernel=linear, gamma=2.22223, C=89 ...  
 [CV] ... kernel=rbf, gamma=3.33334, C=34, total= 3.3s  
 [CV] kernel=linear, gamma=2.22223, C=89 ...  
 [CV] ... kernel=rbf, gamma=1.11112, C=100, total= 3.2s  
 [CV] kernel=linear, gamma=2.22223, C=89 ...  
 [CV] ... kernel=rbf, gamma=1.11112, C=100, total= 3.3s  
 [CV] kernel=rbf, gamma=10.0, C=34 ...  
 [CV] ... kernel=rbf, gamma=10.0, C=34, total= 3.0s  
 [CV] kernel=rbf, gamma=10.0, C=34 ...  
 [CV] ... kernel=rbf, gamma=1.11112, C=100, total= 3.6s  
 [CV] kernel=rbf, gamma=10.0, C=34 ...  
 [CV] ... kernel=linear, gamma=7.77778, C=1, total= 15.1s  
 [CV] kernel=rbf, gamma=3.33334, C=100 ...  
 [CV] ... kernel=rbf, gamma=10.0, C=34, total= 3.1s  
 [CV] kernel=rbf, gamma=3.33334, C=100 ...  
 [CV] ... kernel=rbf, gamma=10.0, C=34, total= 2.9s  
 [CV] kernel=rbf, gamma=3.33334, C=100 ...  
 [CV] ... kernel=rbf, gamma=3.33334, C=100, total= 3.3s  
 [CV] kernel=linear, gamma=8.88889, C=89 ...  
 [CV] ... kernel=rbf, gamma=3.33334, C=100, total= 3.4s  
 [CV] kernel=linear, gamma=8.88889, C=89 ...  
 [CV] ... kernel=rbf, gamma=3.33334, C=100, total= 3.4s  
 [CV] kernel=linear, gamma=8.88889, C=89 ...  
 [CV] ... kernel=linear, gamma=7.77778, C=1, total= 32.7s  
 [CV] kernel=rbf, gamma=8.88889, C=23 ...  
 [CV] ... kernel=rbf, gamma=8.88889, C=23, total= 3.2s  
 [CV] kernel=rbf, gamma=8.88889, C=23 ...  
 [CV] ... kernel=rbf, gamma=8.88889, C=23, total= 2.9s  
 [CV] kernel=rbf, gamma=8.88889, C=23 ...  
 [CV] ... kernel=rbf, gamma=8.88889, C=23, total= 2.9s  
 [CV] kernel=rbf, gamma=6.66667, C=78 ...

[CV] ... kernel=linear, gamma=7.77778, C=1, total= 44.3s  
 [CV] kernel=rbf, gamma=6.66667, C=78 ...  
 [CV] ... kernel=rbf, gamma=6.66667, C=78, total= 3.7s  
 [CV] kernel=rbf, gamma=6.66667, C=78 ...  
 [CV] ... kernel=rbf, gamma=6.66667, C=78, total= 3.8s  
 [CV] kernel=rbf, gamma=1.11112, C=45 ...  
 [CV] ... kernel=rbf, gamma=6.66667, C=78, total= 3.6s  
 [CV] kernel=rbf, gamma=1.11112, C=45 ...  
 [CV] ... kernel=rbf, gamma=1.11112, C=45, total= 3.8s  
 [CV] kernel=rbf, gamma=1.11112, C=45 ...  
 [CV] ... kernel=rbf, gamma=1.11112, C=45, total= 3.6s  
 [CV] kernel=rbf, gamma=4.44445, C=34 ...  
 [CV] ... kernel=rbf, gamma=1.11112, C=45, total= 3.5s  
 [CV] kernel=rbf, gamma=4.44445, C=34 ...  
 [CV] ... kernel=rbf, gamma=4.44445, C=34, total= 3.3s  
 [CV] kernel=rbf, gamma=4.44445, C=34 ...  
 [CV] ... kernel=rbf, gamma=4.44445, C=34, total= 3.2s  
 [CV] kernel=rbf, gamma=2.22223, C=67 ...  
 [CV] ... kernel=rbf, gamma=4.44445, C=34, total= 3.1s  
 [CV] kernel=rbf, gamma=2.22223, C=67 ...  
 [CV] ... kernel=rbf, gamma=2.22223, C=67, total= 3.5s  
 [CV] kernel=rbf, gamma=2.22223, C=67 ...  
 [CV] ... kernel=rbf, gamma=2.22223, C=67, total= 3.6s  
 [CV] kernel=rbf, gamma=8.88889, C=1 ...  
 [CV] ... kernel=rbf, gamma=2.22223, C=67, total= 3.5s  
 [CV] kernel=rbf, gamma=8.88889, C=1 ...  
 [CV] ... kernel=rbf, gamma=8.88889, C=1, total= 2.8s  
 [CV] kernel=rbf, gamma=8.88889, C=1 ...  
 [CV] ... kernel=rbf, gamma=8.88889, C=1, total= 2.7s  
 [CV] kernel=rbf, gamma=4.44445, C=78 ...  
 [CV] ... kernel=rbf, gamma=8.88889, C=1, total= 2.6s  
 [CV] kernel=rbf, gamma=4.44445, C=78 ...  
 [CV] ... kernel=rbf, gamma=4.44445, C=78, total= 3.3s  
 [CV] kernel=rbf, gamma=4.44445, C=78 ...  
 [CV] ... kernel=rbf, gamma=4.44445, C=78, total= 3.3s  
 [CV] kernel=linear, gamma=6.66667, C=45 ...  
 [CV] ... kernel=rbf, gamma=4.44445, C=78, total= 3.3s  
 [CV] kernel=linear, gamma=6.66667, C=45 ...  
 [CV] ... kernel=linear, gamma=4.44445, C=34, total= 1.8min  
 [CV] kernel=linear, gamma=6.66667, C=45 ...  
 [CV] ... kernel=linear, gamma=2.22223, C=23, total= 2.1min  
 [CV] kernel=linear, gamma=2.22223, C=34 ...  
 [CV] ... kernel=linear, gamma=2.22223, C=23, total= 3.0min  
 [CV] kernel=linear, gamma=2.22223, C=34 ...  
 [CV] ... kernel=linear, gamma=7.77778, C=56, total= 3.0min  
 [CV] kernel=linear, gamma=2.22223, C=34 ...  
 [CV] ... kernel=linear, gamma=7.77778, C=45, total= 3.1min  
 [CV] ... kernel=linear, gamma=2.22223, C=23, total= 3.1min



```
[CV] ... kernel=linear, gamma=4.44445, C=34, total= 3.9min
[CV] ... kernel=linear, gamma=7.77778, C=45, total= 4.2min
[CV] ... kernel=linear, gamma=4.44445, C=34, total= 4.3min
[CV] ... kernel=linear, gamma=2.22223, C=34, total= 1.5min
[CV] ... kernel=linear, gamma=2.22223, C=34, total= 2.4min
[CV] ... kernel=linear, gamma=8.88889, C=89, total= 4.3min
[CV] ... kernel=linear, gamma=2.22223, C=89, total= 4.5min
[CV] ... kernel=linear, gamma=6.66667, C=45, total= 2.9min
[CV] ... kernel=linear, gamma=6.66667, C=45, total= 3.5min
[CV] ... kernel=linear, gamma=8.88889, C=89, total= 4.5min
[CV] ... kernel=linear, gamma=7.77778, C=45, total= 4.9min
[CV] ... kernel=linear, gamma=7.77778, C=56, total= 5.0min
[CV] ... kernel=linear, gamma=2.22223, C=34, total= 1.9min
[CV] ... kernel=linear, gamma=2.22223, C=89, total= 4.8min
[CV] ... kernel=linear, gamma=6.66667, C=45, total= 3.7min
[CV] ... kernel=linear, gamma=7.77778, C=56, total= 5.0min
[CV] ... kernel=linear, gamma=8.88889, C=89, total= 4.7min
[CV] ... kernel=linear, gamma=2.22223, C=89, total= 5.0min
```

```
[Parallel(n_jobs=20)]: Done 75 out of 75 | elapsed: 5.1min finished
```

```
-----
NameError                                Traceback (most recent call last)
```

```
<ipython-input-74-e1c77ee7c201> in <module>()
      4
      5 searchScore = metrics.accuracy_score(ytestData, y_svcrand)
----> 6 nonSearchScore = metrics.accuracy_score(ytestData, y_svc)
      7 improvement = (searchScore - nonSearchScore) / nonSearchScore
      8
```

```
NameError: name 'y_svc' is not defined
```

```
In [ ]: # DNN
```

```
hidden_layers = [int(x) for x in np.linspace(start = 10, stop = 500, num = 10)]

activation = ['identity', 'logistic', 'relu', 'tanh']

solver = ['lbfgs', 'adam', 'sgd']

random_grid = {'hidden_layers': hidden_layers,
```

```
'activation': activation,
'solver': solver}
```

```
dnn_random = RandomizedSearchCV(estimator = dnn, param_distributions = random_grid, n_
                               verbose=2, random_state=42, n_jobs = 20)
```

## 10 ROC curves for HOF


```
In [61]: rocHOF, ((ax1, ax2), (ax3, ax4)) = plt.subplots(2, 2, sharey = True, sharex = True)
```

```
ax1.plot(fprSVC, tprSVC, label = 'ROC curve')
ax1.plot([0, 1], [0, 1], linestyle = '--', label = 'Reference line')
ax1.set_title("SVC: %.2f" % roc_aucSVC, size = 21, x = .485, ha = 'center')
```

```
ax3.plot(fprDNN, tprDNN)
ax3.plot([0, 1], [0, 1], linestyle = '--')
ax3.set_title("DNN: %.2f" % roc_aucDNN, size = 21, x = .485, ha = 'center')
```

```
rocHOF.text(-0.03, 0.5, "True positive rate", va='center', rotation='vertical', size = 18)
rocHOF.text(0.5, -0.04, "False positive rate", ha = 'center', size = 18)
```

```
rocHOF.savefig('roc-hof.png', dpi = 400, bbox_inches = 'tight')
```



## 11 Predict rookies

```
In [37]: svcPred = svc.predict(rookieFeatures)

        for i, j in zip(svcPred, rookieNames):
            print(i, j)
```

```
0 Ben Simmons
0 Lonzo Ball
0 Jayson Tatum
0 Josh Jackson
0 De'Aaron Fox
0 Lauri Markkanen
0 Frank Ntilikina
0 Dennis Smith
0 Zach Collins
0 Malik Monk
0 Luke Kennard
0 Donovan Mitchell
0 Bam Adebayo
0 Justin Jackson
0 T.J. Leaf
0 John Collins
0 Terrance Ferguson
0 Jarrett Allen
0 OG Anunoby
0 Kyle Kuzma
0 Josh Hart
0 Wesley Iwundu
0 Frank Mason
0 Semi Ojeleye
0 Jordan Bell
0 Dwayne Bacon
0 Tyler Dorsey
0 Dillon Brooks
0 Sterling Brown
0 Sindarius Thornwell
```

```
In [39]: dnnPred = dnn.predict(rookieFeatures)

        for i, j in zip(dnnPred, rookieNames):
            print(i, j)
```

```
0 Ben Simmons
0 Lonzo Ball
```

```
0 Jayson Tatum
0 Josh Jackson
0 De'Aaron Fox
0 Lauri Markkanen
0 Frank Ntilikina
0 Dennis Smith
0 Zach Collins
0 Malik Monk
0 Luke Kennard
0 Donovan Mitchell
0 Bam Adebayo
0 Justin Jackson
0 T.J. Leaf
0 John Collins
0 Terrance Ferguson
0 Jarrett Allen
0 OG Anunoby
0 Kyle Kuzma
0 Josh Hart
0 Wesley Iwundu
0 Frank Mason
0 Semi Ojeleye
0 Jordan Bell
0 Dwayne Bacon
0 Tyler Dorsey
0 Dillon Brooks
0 Sterling Brown
0 Sindarius Thornwell
```

## 12 Create data splits for All Star models

```
In [40]: trainData, testData = train_test_split(pastDf, test_size = 0.25, random_state = 0)

xtrainData = trainData[['G', 'MPG', 'FG/G', 'FGA/G', '2P%', '3P%', 'FT%', 'TRB/G', 'AST',
                        'TS%', '3PAr', 'FTr']]
ytrainData = trainData[['All Star']]

xtestData = testData[['G', 'MPG', 'FG/G', 'FGA/G', '2P%', '3P%', 'FT%', 'TRB/G', 'AST',
                      'TS%', '3PAr', 'FTr']]
ytestData = testData[['All Star']]
```

## 13 Create models and confusion matrices for All Star models

```
In [71]: svcModel = SVC(kernel='rbf', gamma=1e-4, C=100, probability = True)
        svcModel.fit(xtrainData, ytrainData.values.ravel())
```

```

Ysvc = svcModel.predict(xtestData)

print("Accuracy score: %.3f" % metrics.accuracy_score(ytestData, Ysvc))

proba = svc.predict_proba(xtestData)
print("Log loss: %.3f" % metrics.log_loss(ytestData, proba))

posProb = proba[:, 1]
print("Area under ROC curve: %.3f" % metrics.roc_auc_score(ytestData, posProb))

fprSVC, tprSVC, thresholdSVC = metrics.roc_curve(ytestData, posProb)
roc_aucSVC = metrics.auc(fprSVC, tprSVC)

cvScoreSVC = cross_val_score(svcModel, xtestData, ytestData.values.ravel(), cv = 3, scoring='accuracy')
print("Accuracy (cross validation score): %0.2f (+/- %0.2f)" % (cvScoreSVC.mean(), cvScoreSVC.std()))

Accuracy score: 0.927
Log loss: 0.231
Area under ROC curve: 0.729
Accuracy (cross validation score): 0.92 (+/- 0.00)

In [72]: cconfusionmatrix = metrics.confusion_matrix(ytestData, y_svc)
         svcHofAS, ax = plt.subplots()

         sns.heatmap(confusionmatrix, annot=True, ax = ax, linewidth = 2)

         ax.set_xlabel("Predicted")
         ax.set_ylabel("Actual")


         labels = ["Not All Star", "All Star"]
         ax.set_xticklabels(labels)
         ax.set_yticklabels(labels)

         svcHofAS.suptitle("SVC Confusion Matrix", weight = 'bold', size = 18, y = 1.04, x = 0.5)

         svcHofAS.savefig('svc-hof-as.png', dpi = 400, bbox_inches = 'tight')

```

## SVC Confusion Matrix



```
In [46]: dnn = MLPClassifier(
          solver='lbfgs',
          hidden_layer_sizes=100,
          max_iter=10000,
          shuffle=False,
          random_state=0,
          activation='identity')

dnn.fit(xtrainData, ytrainData.values.ravel())

y_dnn = dnn.predict(xtestData)

print("Accuracy score: %.3f" % metrics.accuracy_score(ytestData, y_dnn))

proba = dnn.predict_proba(xtestData)
print("Log loss: %.3f" % metrics.log_loss(ytestData, proba))

posProb = proba[:, 1]
print("Area under ROC curve: %.3f" % metrics.roc_auc_score(ytestData, posProb))

fprDNN, tprDNN, thresholdDNN = metrics.roc_curve(ytestData, posProb)
```

```

roc_aucDNN = metrics.auc(fprDNN, tprDNN)

cvScoreDNN = cross_val_score(dnn, xtestData, ytestData.values.ravel(), cv = 3, scoring='roc_auc')
print("Accuracy (cross validation score): %0.2f (+/- %0.2f)" % (cvScoreDNN.mean(), cvScoreDNN.std()))

Accuracy score: 0.930
Log loss: 0.208
Area under ROC curve: 0.846
Accuracy (cross validation score): 0.93 (+/- 0.01)

In [73]: confusionmatrix = metrics.confusion_matrix(ytestData, y_dnn)

dnnHofAS, ax = plt.subplots()

sns.heatmap(confusionmatrix, annot=True, ax = ax, linewidth = 2)

ax.set_xlabel("Predicted")
ax.set_ylabel("Actual")

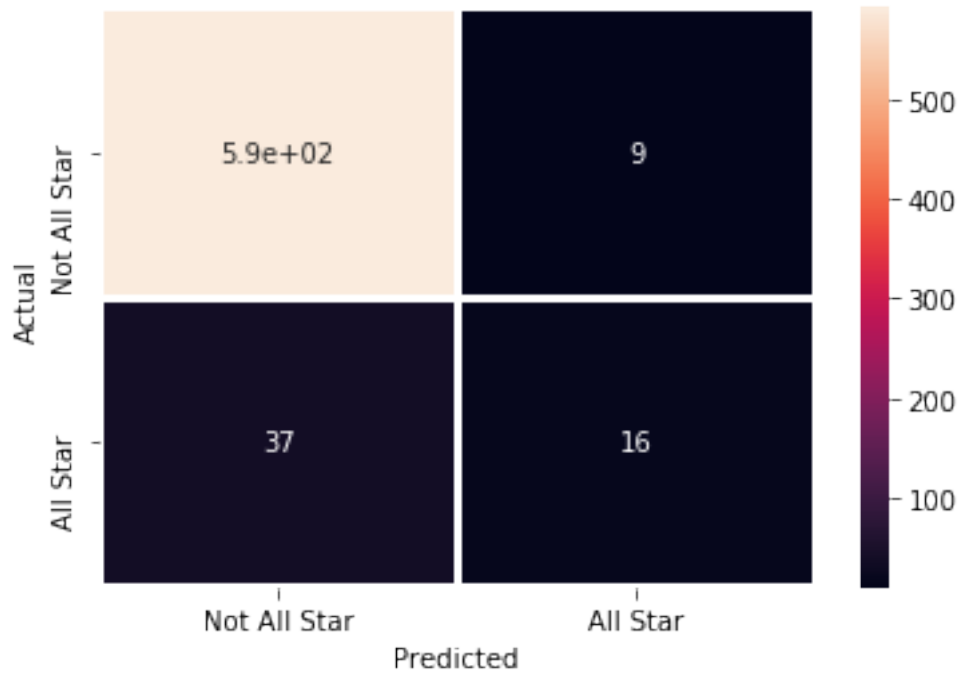
labels = ["Not All Star", "All Star"]
ax.set_xticklabels(labels)
ax.set_yticklabels(labels)

dnnHofAS.suptitle("DNN Confusion Matrix", weight = 'bold', size = 18, y = 1.04, x = .4)

dnnHofAS.savefig('dnn-hof-as.png', dpi = 400, bbox_inches = 'tight')

```

## DNN Confusion Matrix



## 14 Randomized search CV

In [48]: # SVC

```
C = [int(x) for x in np.linspace(start = 1, stop = 100, num = 10)]

kernel = ['rbf', 'linear']

gamma = [float(x) for x in np.linspace(start = 1e-5, stop = 10, num = 10)]

random_grid = {'C': C,
               'kernel': kernel,
               'gamma': gamma}

svc_random = RandomizedSearchCV(estimator = svc, param_distributions = random_grid, n
                                verbose=2, random_state=42, n_jobs = 20)
```

In [49]: `svc_random.fit(xtrainData, ytrainData.values.ravel())`

```
y_svcrand = svc_random.predict(xtestData)
```



```

searchScore = metrics.accuracy_score(ytestData, y_svcrand)
nonSearchScore = metrics.accuracy_score(ytestData, y_svc)
improvement = (searchScore - nonSearchScore) / nonSearchScore

```

```

print("Percent improvement: %.5f" % improvement)

```

Fitting 3 folds for each of 25 candidates, totalling 75 fits

```

[CV] kernel=linear, gamma=7.77778, C=45 ...
[CV] kernel=linear, gamma=7.77778, C=45 ...
[CV] kernel=linear, gamma=7.77778, C=45 ...
[CV] kernel=linear, gamma=7.77778, C=1 ...
[CV] kernel=linear, gamma=7.77778, C=1 ...
[CV] kernel=linear, gamma=7.77778, C=1 ...
[CV] kernel=rbf, gamma=5.55556, C=12 ...
[CV] kernel=rbf, gamma=5.55556, C=12 ...
[CV] kernel=rbf, gamma=5.55556, C=12 ...
[CV] kernel=rbf, gamma=10.0, C=78 ...
[CV] kernel=rbf, gamma=10.0, C=78 ...
[CV] kernel=rbf, gamma=10.0, C=78 ...
[CV] kernel=rbf, gamma=4.44445, C=67 ...
[CV] kernel=rbf, gamma=4.44445, C=67 ...
[CV] kernel=rbf, gamma=4.44445, C=67 ...
[CV] kernel=linear, gamma=7.77778, C=56 ...
[CV] kernel=linear, gamma=7.77778, C=56 ...
[CV] kernel=linear, gamma=7.77778, C=56 ...
[CV] kernel=linear, gamma=4.44445, C=34 ...
[CV] kernel=linear, gamma=4.44445, C=34 ...
[CV] ... kernel=rbf, gamma=10.0, C=78, total= 3.1s
[CV] kernel=linear, gamma=4.44445, C=34 ...
[CV] ... kernel=rbf, gamma=10.0, C=78, total= 3.2s
[CV] kernel=rbf, gamma=5.55556, C=89 ...
[CV] ... kernel=rbf, gamma=5.55556, C=12, total= 3.3s
[CV] kernel=rbf, gamma=5.55556, C=89 ...
[CV] ... kernel=rbf, gamma=4.44445, C=67, total= 3.3s
[CV] ... kernel=rbf, gamma=10.0, C=78, total= 3.3s
[CV] ... kernel=rbf, gamma=5.55556, C=12, total= 3.4s

```

```

[Parallel(n_jobs=20)]: Done 1 tasks | elapsed: 3.6s

```

```

[CV] kernel=rbf, gamma=7.77778, C=89 ...
[CV] kernel=rbf, gamma=7.77778, C=89 ...
[CV] kernel=rbf, gamma=5.55556, C=89 ...
[CV] ... kernel=rbf, gamma=5.55556, C=12, total= 3.4s
[CV] kernel=rbf, gamma=7.77778, C=89 ...
[CV] ... kernel=rbf, gamma=4.44445, C=67, total= 3.4s
[CV] kernel=linear, gamma=2.22223, C=23 ...

```

[CV] ... kernel=rbf, gamma=4.44445, C=67, total= 3.5s  
 [CV] kernel=linear, gamma=2.22223, C=23 ...  
 [CV] ... kernel=rbf, gamma=7.77778, C=89, total= 3.1s  
 [CV] kernel=linear, gamma=2.22223, C=23 ...  
 [CV] ... kernel=rbf, gamma=5.55556, C=89, total= 3.3s  
 [CV] kernel=rbf, gamma=3.33334, C=34 ...  
 [CV] ... kernel=rbf, gamma=7.77778, C=89, total= 3.2s  
 [CV] kernel=rbf, gamma=3.33334, C=34 ...  
 [CV] ... kernel=rbf, gamma=7.77778, C=89, total= 3.3s  
 [CV] ... kernel=rbf, gamma=5.55556, C=89, total= 3.4s  
 [CV] kernel=rbf, gamma=3.33334, C=34 ...  
 [CV] kernel=rbf, gamma=1.11112, C=100 ...  
 [CV] ... kernel=rbf, gamma=5.55556, C=89, total= 3.3s  
 [CV] kernel=rbf, gamma=1.11112, C=100 ...  
 [CV] ... kernel=linear, gamma=7.77778, C=1, total= 10.0s  
 [CV] kernel=rbf, gamma=1.11112, C=100 ...  
 [CV] ... kernel=rbf, gamma=3.33334, C=34, total= 3.3s  
 [CV] kernel=linear, gamma=2.22223, C=89 ...  
 [CV] ... kernel=rbf, gamma=3.33334, C=34, total= 3.3s  
 [CV] ... kernel=rbf, gamma=3.33334, C=34, total= 3.4s  
 [CV] kernel=linear, gamma=2.22223, C=89 ...  
 [CV] kernel=linear, gamma=2.22223, C=89 ...  
 [CV] ... kernel=rbf, gamma=1.11112, C=100, total= 3.3s  
 [CV] kernel=rbf, gamma=10.0, C=34 ...  
 [CV] ... kernel=rbf, gamma=1.11112, C=100, total= 3.3s  
 [CV] kernel=rbf, gamma=10.0, C=34 ...  
 [CV] ... kernel=linear, gamma=7.77778, C=1, total= 11.9s  
 [CV] kernel=rbf, gamma=10.0, C=34 ...  
 [CV] ... kernel=rbf, gamma=1.11112, C=100, total= 3.6s  
 [CV] kernel=rbf, gamma=3.33334, C=100 ...  
 [CV] ... kernel=rbf, gamma=10.0, C=34, total= 3.1s  
 [CV] kernel=rbf, gamma=3.33334, C=100 ...  
 [CV] ... kernel=rbf, gamma=10.0, C=34, total= 3.3s  
 [CV] kernel=rbf, gamma=3.33334, C=100 ...  
 [CV] ... kernel=rbf, gamma=10.0, C=34, total= 3.1s  
 [CV] kernel=linear, gamma=8.88889, C=89 ...  
 [CV] ... kernel=rbf, gamma=3.33334, C=100, total= 3.3s  
 [CV] kernel=linear, gamma=8.88889, C=89 ...  
 [CV] ... kernel=rbf, gamma=3.33334, C=100, total= 3.3s  
 [CV] kernel=linear, gamma=8.88889, C=89 ...  
 [CV] ... kernel=rbf, gamma=3.33334, C=100, total= 3.2s  
 [CV] kernel=rbf, gamma=8.88889, C=23 ...  
 [CV] ... kernel=rbf, gamma=8.88889, C=23, total= 2.9s  
 [CV] kernel=rbf, gamma=8.88889, C=23 ...  
 [CV] ... kernel=rbf, gamma=8.88889, C=23, total= 2.7s  
 [CV] kernel=rbf, gamma=8.88889, C=23 ...  
 [CV] ... kernel=rbf, gamma=8.88889, C=23, total= 2.8s  
 [CV] kernel=rbf, gamma=6.66667, C=78 ...

[CV] ... kernel=rbf, gamma=6.66667, C=78, total= 2.9s  
 [CV] kernel=rbf, gamma=6.66667, C=78 ...  
 [CV] ... kernel=linear, gamma=7.77778, C=1, total= 33.5s  
 [CV] kernel=rbf, gamma=6.66667, C=78 ...  
 [CV] ... kernel=rbf, gamma=6.66667, C=78, total= 3.1s  
 [CV] kernel=rbf, gamma=1.11112, C=45 ...  
 [CV] ... kernel=rbf, gamma=6.66667, C=78, total= 2.9s  
 [CV] kernel=rbf, gamma=1.11112, C=45 ...  
 [CV] ... kernel=rbf, gamma=1.11112, C=45, total= 3.1s  
 [CV] kernel=rbf, gamma=1.11112, C=45 ...  
 [CV] ... kernel=rbf, gamma=1.11112, C=45, total= 3.1s  
 [CV] kernel=rbf, gamma=4.44445, C=34 ...  
 [CV] ... kernel=rbf, gamma=1.11112, C=45, total= 3.2s  
 [CV] kernel=rbf, gamma=4.44445, C=34 ...  
 [CV] ... kernel=rbf, gamma=4.44445, C=34, total= 2.9s  
 [CV] kernel=rbf, gamma=4.44445, C=34 ...  
 [CV] ... kernel=rbf, gamma=4.44445, C=34, total= 3.0s  
 [CV] kernel=rbf, gamma=2.22223, C=67 ...  
 [CV] ... kernel=rbf, gamma=4.44445, C=34, total= 3.0s  
 [CV] kernel=rbf, gamma=2.22223, C=67 ...  
 [CV] ... kernel=rbf, gamma=2.22223, C=67, total= 3.2s  
 [CV] kernel=rbf, gamma=2.22223, C=67 ...  
 [CV] ... kernel=rbf, gamma=2.22223, C=67, total= 3.1s  
 [CV] kernel=rbf, gamma=8.88889, C=1 ...  
 [CV] ... kernel=rbf, gamma=2.22223, C=67, total= 3.1s  
 [CV] kernel=rbf, gamma=8.88889, C=1 ...  
 [CV] ... kernel=rbf, gamma=8.88889, C=1, total= 2.7s  
 [CV] kernel=rbf, gamma=8.88889, C=1 ...  
 [CV] ... kernel=rbf, gamma=8.88889, C=1, total= 2.7s  
 [CV] kernel=rbf, gamma=4.44445, C=78 ...  
 [CV] ... kernel=rbf, gamma=8.88889, C=1, total= 2.7s  
 [CV] kernel=rbf, gamma=4.44445, C=78 ...  
 [CV] ... kernel=rbf, gamma=4.44445, C=78, total= 2.9s  
 [CV] kernel=rbf, gamma=4.44445, C=78 ...  
 [CV] ... kernel=rbf, gamma=4.44445, C=78, total= 3.0s  
 [CV] kernel=linear, gamma=6.66667, C=45 ...  
 [CV] ... kernel=rbf, gamma=4.44445, C=78, total= 3.1s  
 [CV] kernel=linear, gamma=6.66667, C=45 ...  
 [CV] ... kernel=linear, gamma=2.22223, C=23, total= 2.6min  
 [CV] kernel=linear, gamma=6.66667, C=45 ...  
 [CV] ... kernel=linear, gamma=4.44445, C=34, total= 3.8min  
 [CV] kernel=linear, gamma=2.22223, C=34 ...  
 [CV] ... kernel=linear, gamma=2.22223, C=23, total= 3.8min  
 [CV] kernel=linear, gamma=2.22223, C=34 ...  
 [CV] ... kernel=linear, gamma=4.44445, C=34, total= 4.3min  
 [CV] kernel=linear, gamma=2.22223, C=34 ...  
 [CV] ... kernel=linear, gamma=7.77778, C=45, total= 4.8min  
 [CV] ... kernel=linear, gamma=2.22223, C=23, total= 4.9min

```
[CV] ... kernel=linear, gamma=7.77778, C=45, total= 5.1min
[CV] ... kernel=linear, gamma=4.44445, C=34, total= 5.6min
[CV] ... kernel=linear, gamma=7.77778, C=56, total= 6.2min
[CV] ... kernel=linear, gamma=7.77778, C=56, total= 6.3min
[CV] ... kernel=linear, gamma=6.66667, C=45, total= 5.3min
[CV] ... kernel=linear, gamma=2.22223, C=34, total= 2.8min
[CV] ... kernel=linear, gamma=6.66667, C=45, total= 4.4min
[CV] ... kernel=linear, gamma=6.66667, C=45, total= 6.3min
[CV] ... kernel=linear, gamma=2.22223, C=34, total= 3.0min
[CV] ... kernel=linear, gamma=8.88889, C=89, total= 7.2min
[CV] ... kernel=linear, gamma=2.22223, C=89, total= 7.4min
[CV] ... kernel=linear, gamma=7.77778, C=45, total= 7.6min
[CV] ... kernel=linear, gamma=2.22223, C=89, total= 7.5min
[CV] ... kernel=linear, gamma=2.22223, C=34, total= 3.9min
[CV] ... kernel=linear, gamma=7.77778, C=56, total= 7.8min
[CV] ... kernel=linear, gamma=8.88889, C=89, total= 7.6min
[CV] ... kernel=linear, gamma=2.22223, C=89, total= 8.0min
[CV] ... kernel=linear, gamma=8.88889, C=89, total= 7.9min
```

```
[Parallel(n_jobs=20)]: Done 75 out of 75 | elapsed: 8.2min finished
```

Percent improvement: 0.00000

In [52]: # DNN

```
hidden_layers = [int(x) for x in np.linspace(start = 10, stop = 500, num = 10)]

activation = ['identity', 'logistic', 'relu', 'tanh']

solver = ['lbfgs', 'adam', 'sgd']

random_grid = {'hidden_layers': hidden_layers,
               'activation': activation,
               'solver': solver}

dnn_random = RandomizedSearchCV(estimator = dnn, param_distributions = random_grid, n
                                verbose=2, random_state=42, n_jobs = 20)
```

## 15 ROC curves for All Star

In [74]: rocAS, ((ax1, ax2), (ax3, ax4)) = plt.subplots(2, 2, sharey = True, sharex = True)

```
ax1.plot(fprSVC, tprSVC, label = 'ROC curve')
ax1.plot([0, 1], [0, 1], linestyle = '--', label = 'Reference line')
ax1.set_title("SVC: %.2f" % roc_aucSVC, size = 21, x = .485, ha = 'center')
```

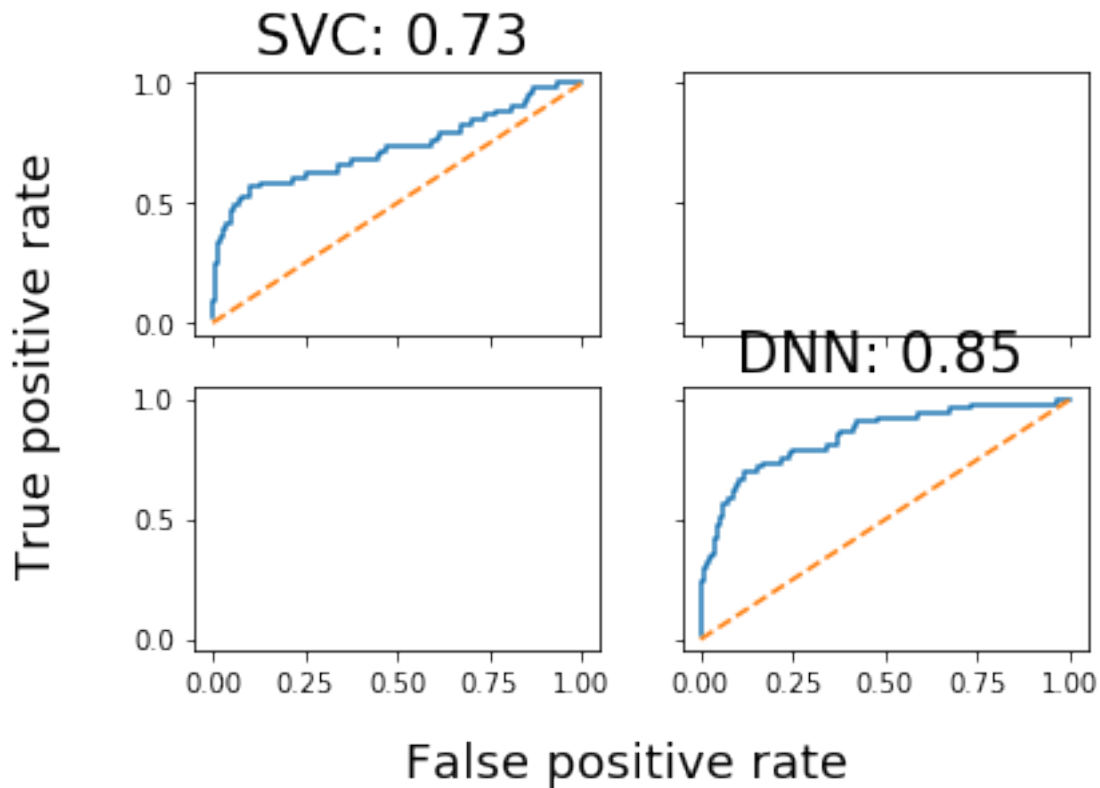
```

ax4.plot(fprDNN, tprDNN)
ax4.plot([0, 1], [0, 1], linestyle = '--')
ax4.set_title("DNN: %.2f" % roc_aucDNN, size = 21, x = .485, ha = 'center')

rocAS.text(-0.03, 0.5, "True positive rate", va='center', rotation='vertical', size =
rocAS.text(0.5, -0.04, "False positive rate", ha = 'center', size = 18)

rocAS.savefig('roc-as.png', dpi = 400, bbox_inches = 'tight')

```



## 16 Predict rookies

```

In [56]: svcPred = svc.predict(rookieFeatures)

for i, j in zip(svcPred, rookieNames):
    print(i, j)

```

```

1 Ben Simmons
0 Lonzo Ball
0 Jayson Tatum
0 Josh Jackson

```

```

0 De'Aaron Fox
0 Lauri Markkanen
0 Frank Ntilikina
0 Dennis Smith
0 Zach Collins
0 Malik Monk
0 Luke Kennard
0 Donovan Mitchell
0 Bam Adebayo
0 Justin Jackson
0 T.J. Leaf
0 John Collins
0 Terrance Ferguson
0 Jarrett Allen
0 OG Anunoby
0 Kyle Kuzma
0 Josh Hart
0 Wesley Iwundu
0 Frank Mason
0 Semi Ojeleye
0 Jordan Bell
0 Dwayne Bacon
0 Tyler Dorsey
0 Dillon Brooks
0 Sterling Brown
0 Sindarius Thornwell

```

```

In [58]: dnnPred = dnn.predict(rookieFeatures)

        for i, j in zip(dnnPred, rookieNames):
            print(i, j)

```

```

1 Ben Simmons
0 Lonzo Ball
0 Jayson Tatum
0 Josh Jackson
0 De'Aaron Fox
0 Lauri Markkanen
0 Frank Ntilikina
0 Dennis Smith
0 Zach Collins
0 Malik Monk
0 Luke Kennard
1 Donovan Mitchell
0 Bam Adebayo
0 Justin Jackson
0 T.J. Leaf

```

0 John Collins  
0 Terrance Ferguson  
0 Jarrett Allen  
0 OG Anunoby  
0 Kyle Kuzma  
0 Josh Hart  
0 Wesley Iwundu  
0 Frank Mason  
0 Semi Ojeleye  
0 Jordan Bell  
0 Dwayne Bacon  
0 Tyler Dorsey  
0 Dillon Brooks  
0 Sterling Brown  
0 Sindarius Thornwell